

## Deepdive 环境搭建及操作

### 一、虚拟机环境配置

查看 python 版本

```
frankstein@ubuntu:~$ python --version
Python 2.7.12
```

安装 openjdk-8-jdk

查看 java 版本，验证是否安装成功

```
frankstein@ubuntu:~$ java -version
openjdk version "1.8.0_292"
OpenJDK Runtime Environment (build 1.8.0_292-8u292-b10-0ubuntu1~16.04.1-b10)
OpenJDK 64-Bit Server VM (build 25.292-b10, mixed mode)
```

### 二、安装 CNDeepdive

为 install.sh 和 nlp\_setup.sh 文件增加可执行权限

```
frankstein@ubuntu:~/CNdeepdive/CNdeepdive$ ls
deepdive-v0.8-STABLE-Darwin.tar.gz
deepdive-v0.8-STABLE-Linux.tar.gz
dependency
install.sh
nlp_setup.sh
readme.md
transaction
tutorial
frankstein@ubuntu:~/CNdeepdive/CNdeepdive$ chmod +x install.sh
frankstein@ubuntu:~/CNdeepdive/CNdeepdive$ chmod +x nlp_setup.sh
```

打开 install.sh 文件，修改第 193 行内容

```
mkdir -p "$PREFIX"
tar xvf "$tarball" -C "$PREFIX"
) || return $?
```

运行 install.sh 文件，输入 1 安装 deepdive

```
+ sudo localedef -i en_US -f UTF-8 en_US.UTF-8
## Finished installation for _deepdive_runtime_deps
## Finished installation for deepdive
# Install what (enter to repeat options, a to see all, q to quit, or a number)? q
frankstein@ubuntu:~/CNdeepdive/CNdeepdive$
```

完成后需要配置 CNdeepdive 的环境变量 添加下面的语句

```
export PATH="$~/local/bin:$PATH"
# enable programmable completion features (you don't need to enable
# this, if it's already enabled in /etc/bash.bashrc and /etc/profile
# sources /etc/bash.bashrc).
```

运行下面的语句让改动生效

在控制台输入 deepdive，出现下面的输出证明配置成功

```
frankstein@ubuntu:~/CNdeepdive/CNdeepdive$ source ~/.bashrc
frankstein@ubuntu:~/CNdeepdive/CNdeepdive$ deepdive
# DeepDive (v0.8.0-79-g28a58de)
# Usage: deepdive COMMAND [ARG]...

# To enable bash completion for deepdive commands, run:
source $(deepdive whereis installed etc/deepdive_bash_completion.sh)

No COMMAND given

# Available COMMANDs are:
deepdive check      # Checks errors in compiled DeepDive app
deepdive compile    # Compiles DeepDive source code into executables under run/
deepdive compute    # Runs a UDF using a computer against the database
deepdive create      # Creates a table/view in the database
deepdive db          # Exposes lower level database operations provided by the driver
deepdive do          # Runs necessary processes to get something done
deepdive done        # Checks whether given targets are all done
```

### 三、安装 postgresql

运行 install.sh 文件，输入 6 安装 postgresql

```
+ sudo service postgresql restart
++ rm -f /tmp/pg_hba.conf.nh0JqJW
## Finished installation for postgres
# Install what (enter to repeat options, a to see all, q to quit, or a number)? q
frankstein@ubuntu:~/CNdeepdive/CNdeepdive$ sudo su - postgres
postgres@ubuntu:~$
```

### 四、安装 nlp

运行 nlp\_setup.sh 文件

```
frankstein@ubuntu:~/CNdeepdive/CNdeepdive$ ./nlp_setup.sh
Dependency Already Installed.
```

### 五、Smoke Example

进入 postgresql 控制台

```
frankstein@ubuntu:~/CNdeepdive/CNdeepdive$ sudo su - postgres
[sudo] password for frankstein:
postgres@ubuntu:~$ psql
psql (9.5.25)
Type "help" for help.

postgres=#
```

创建用户 设置用户权限

```
postgres=# CREATE USER ke WITH PASSWORD 'frankstein';
CREATE ROLE
postgres=# ALTER ROLE ke CREATEROLE CREATEDB;
ALTER ROLE
```

创建 smoke\_test 数据库，并指定所有者

退出 postgresql 控制台

```
postgres=# CREATE DATABASE smoke_test OWNER ke;
CREATE DATABASE
postgres=# \q
postgres@ubuntu:~$
```

查看是否创建成功

```
postgres@ubuntu:~$ psql -U ke -d smoke_test -h 127.0.0.1 -p 5432
psql (9.5.25)
SSL connection (protocol: TLSv1.2, cipher: ECDHE-RSA-AES256-GCM-SHA3
84, bits: 256, compression: off)
Type "help" for help.

smoke_test=>
```

进入 smoke 文件夹后，修改用户权限

建立数据库配置文件

```
frankstein@ubuntu:~/CNdeepdive/CNdeepdive$ cd smoke
frankstein@ubuntu:~/CNdeepdive/CNdeepdive/smoke$ chmod -R 777 ./
frankstein@ubuntu:~/CNdeepdive/CNdeepdive/smoke$ echo "postgresql://ke@localhost:5432/smoke_test">db.url
frankstein@ubuntu:~/CNdeepdive/CNdeepdive/smoke$
```

编译



```

frankstein@ubuntu:~/CNdeepdive/CNdeepdive/smoke/smoke$ deepdive compile
2025-05-27 20:53:04.538796 'run/LATEST.COMPILE' -> '20250527/205304.452027345'
2025-05-27 20:53:04.539198 'run/RUNNING.COMPILE' -> '20250527/205304.452027345'
2025-05-27 20:53:04.539425 Parsing DeepDive application (/home/frankstein/CNdeepdive/CNdeepdive/smoke/smoke)
to generate:
2025-05-27 20:53:04.539534 run/compiled/schema.json
2025-05-27 20:53:04.539662 from app.ddlog
2025-05-27 20:53:06.274008 run/compiled/deepdive.conf
2025-05-27 20:53:06.274113 from app.ddlog
2025-05-27 20:53:07.068709 from deepdive.conf
2025-05-27 20:53:07.074563 run/compiled/deepdive.conf.json
2025-05-27 20:53:07.733721 Performing sanity checks on run/compiled/deepdive.conf.json:
2025-05-27 20:53:07.809170 checking if input_extractors well defined
2025-05-27 20:53:07.809266 checking if input_schema wellformed
2025-05-27 20:53:07.809683 Normalizing and adding built-in processes to the data flow to compile:
2025-05-27 20:53:07.810316 run/compiled/config-0.00-init_objects.json
2025-05-27 20:53:07.830806 run/compiled/config-0.01-parse_calibration.json
2025-05-27 20:53:07.855720 run/compiled/config-0.01-parse_schema.json
2025-05-27 20:53:07.889410 run/compiled/config-0.51-add_init_app.json
2025-05-27 20:53:07.905459 run/compiled/config-0.52-input_loader.json
2025-05-27 20:53:07.924202 run/compiled/config-1.00-qualified_names.json
2025-05-27 20:53:07.948232 run/compiled/config-1.01-parse_inference_rules.json
2025-05-27 20:53:08.002113 run/compiled/config-2.01-grounding.json
2025-05-27 20:53:08.164106 run/compiled/config-2.02-learning_inference.json
2025-05-27 20:53:08.206816 run/compiled/config-2.03-calibration_plots.json
2025-05-27 20:53:08.230885 run/compiled/config-9.98-ensure_init_app.json
2025-05-27 20:53:08.257888 run/compiled/config-9.99-dependencies.json
2025-05-27 20:53:08.286843 run/compiled/config.json
2025-05-27 20:53:08.291022 Validating run/compiled/config.json:
2025-05-27 20:53:08.403907 checking if compiled_base_relations_have_input_data
2025-05-27 20:53:08.403970 checking if compiled_dependencies_correct
2025-05-27 20:53:08.403981 checking if compiled_input_output_well_defined
2025-05-27 20:53:08.403989 checking if compiled_output_uniquely_defined
2025-05-27 20:53:08.405211 Compiling executable code into:
2025-05-27 20:53:08.405275 run/compiled/code-cmd_extractor.json
2025-05-27 20:53:08.442789 run/compiled/code-dataflow_dot.json
2025-05-27 20:53:08.486321 run/compiled/code-Makefile.json
2025-05-27 20:53:08.537956 run/compiled/code-sql_extractor.json
2025-05-27 20:53:08.561861 run/compiled/code-tsv_extractor.json
2025-05-27 20:53:08.598070 Generating files:
2025-05-27 20:53:08.634853 run/process/init/app/run.sh
2025-05-27 20:53:08.642762 run/process/init/relation/friends/run.sh
2025-05-27 20:53:08.652242 run/process/init/relation/person/run.sh
2025-05-27 20:53:08.660866 run/process/init/relation/person_has_cancer/run.sh
2025-05-27 20:53:08.671763 run/process/init/relation/person_smokes/run.sh
2025-05-27 20:53:08.680263 run/process/grounding/variable_id_partition/run.sh
2025-05-27 20:53:08.687366 run/process/grounding/variable/person_has_cancer/assign_id/run.sh
2025-05-27 20:53:08.694186 run/process/grounding/variable/person_smokes/assign_id/run.sh
2025-05-27 20:53:08.704180 run/process/grounding/variable_holdout/run.sh
2025-05-27 20:53:08.714438 run/process/grounding/variable/person_has_cancer/dump/run.sh
2025-05-27 20:53:08.723398 run/process/grounding/variable/person_smokes/dump/run.sh
2025-05-27 20:53:08.732594 run/process/grounding/factor/inf_imply_person_smokes_person_has_cancer/materiali
ze/run.sh
2025-05-27 20:53:08.739511 run/process/grounding/factor/inf_imply_person_smokes_person_smokes/materialize/r
un.sh
2025-05-27 20:53:08.746957 run/process/grounding/weight_id_partition/run.sh
2025-05-27 20:53:08.756792 run/process/grounding/global_weight_table/run.sh
2025-05-27 20:53:08.765260 run/process/grounding/factor/inf_imply_person_smokes_person_has_cancer/assign_we
ight_id/run.sh
2025-05-27 20:53:08.784084 run/process/grounding/factor/inf_imply_person_smokes_person_smokes/assign_weight
_id/run.sh
2025-05-27 20:53:08.792086 run/process/grounding/factor/inf_imply_person_smokes_person_has_cancer/dump/run.
sh
2025-05-27 20:53:08.800225 run/process/grounding/factor/inf_imply_person_smokes_person_smokes/dump/run.sh

2025-05-27 20:53:08.877530 run/process/model/load_probabilities/run.sh
2025-05-27 20:53:08.886542 run/process/model/calibration/run.sh
2025-05-27 20:53:08.895507 run/dataflow.dot
2025-05-27 20:53:08.902415 run/Makefile
2025-05-27 20:53:08.988513 run/dataflow.svg
2025-05-27 20:53:08.990270 (file:///home/frankstein/CNdeepdive/CNdeepdive/smoke/smoke/run/dataflow.svg)
2025-05-27 20:53:08.993518 'run/compiled' -> '20250527/205304.452027345'

```

运行



```

frankstein@ubuntu:~/CNdeepdive/CNdeepdive/smoke/smoke$ deepdive run
'run/RUNNING' -> '20250527/205405.863438859'
'run/LATEST' -> '20250527/205405.863438859'
2025-05-27 20:54:13.239940 # on ubuntu: deepdive do process/init/app model/calibration-plots
2025-05-27 20:54:13.240139 # run/20250527/205405.863438859/plan.sh
2025-05-27 20:54:13.240194 # execution plan for process/init/app model/calibration-plots
2025-05-27 20:54:13.240236
2025-05-27 20:54:13.240276 ## process/init/app #####
2025-05-27 20:54:13.240316 # Done: N/A
2025-05-27 20:54:13.240361 process/init/app/run.sh
2025-05-27 20:54:13.240399 ++ dirname process/init/app/run.sh
2025-05-27 20:54:13.240436 + cd process/init/app
2025-05-27 20:54:13.240489 + export DEEPDIVE_CURRENT_PROCESS_NAME=process/init/app
2025-05-27 20:54:13.240533 + DEEPDIVE_CURRENT_PROCESS_NAME=process/init/app
2025-05-27 20:54:13.240575 + deepdive db init
2025-05-27 20:54:13.897616 + cd /home/frankstein/CNdeepdive/CNdeepdive/smoke/smoke
2025-05-27 20:54:13.897863 + [[ -r schema.sql ]]
2025-05-27 20:54:13.897906 + [[ -x input/init.sh ]]
2025-05-27 20:54:13.898339 mark_done process/init/app
2025-05-27 20:54:13.940762 #####
2025-05-27 20:54:13.940836
2025-05-27 20:54:13.940848
2025-05-27 20:54:13.940857 ## process/init/relation/person #####
2025-05-27 20:54:13.940865 # Done: N/A
2025-05-27 20:54:13.940872 process/init/relation/person/run.sh
2025-05-27 20:54:13.941198 ++ dirname process/init/relation/person/run.sh
2025-05-27 20:54:13.942804 + cd process/init/relation/person
2025-05-27 20:54:13.943190 + export DEEPDIVE_CURRENT_PROCESS_NAME=process/init/relation/person
2025-05-27 20:54:13.943303 + DEEPDIVE_CURRENT_PROCESS_NAME=process/init/relation/person
2025-05-27 20:54:13.943382 + deepdive create table person
2025-05-27 20:54:14.211267 NOTICE: table "person" does not exist, skipping
2025-05-27 20:54:14.217578 CREATE TABLE

ve/CNdeepdive/smoke/smoke/run/20250527/205405.863438859/calibration/
2025-05-27 20:54:26.264765 mark_done process/model/calibration
2025-05-27 20:54:26.278450 #####
2025-05-27 20:54:26.278697
2025-05-27 20:54:26.280915
2025-05-27 20:54:26.281009 ## model/calibration-plots #####
2025-05-27 20:54:26.281332 # Done: N/A
2025-05-27 20:54:26.281430 mark_done model/calibration-plots
2025-05-27 20:54:26.293731 #####
2025-05-27 20:54:26.293780
2025-05-27 20:54:26.293787
'run/FINISHED' -> '20250527/205405.863438859'

```

运行完成后查看结果

```

frankstein@ubuntu:~/CNdeepdive/CNdeepdive/smoke/smoke$ psql smoke_test
psql (9.5.25)
Type "help" for help.

smoke_test=# \d

```

Schema	Name	Type	Owner
public	dd_factors_inf_imply_person_smokes_person_has_cancer	table	ke
public	dd_factors_inf_imply_person_smokes_person_smokes	table	ke
public	dd_graph_variables_holdout	table	ke
public	dd_graph_variables_observation	table	ke
public	dd_graph_weights	view	ke
public	dd_inference_result_variables	table	ke
public	dd_weights_inf_imply_person_smokes_person_has_cancer	table	ke
public	dd_weights_inf_imply_person_smokes_person_smokes	table	ke
public	friends	table	ke
public	person	table	ke
public	person_has_cancer	table	ke
public	person_has_cancer_label_calibration	view	ke
public	person_has_cancer_label_inference	view	ke
public	person_smokes	table	ke
public	person_smokes_label_calibration	view	ke
public	person_smokes_label_inference	view	ke

```

(16 rows)

```

查看推理结果

```
smoke_test=# select * from person_has_cancer_label_inference;  
person_id | id | label | category | expectation
```

person_id	id	label	category	expectation
1	0		1	0.633
3	2		1	0.628
4	3		1	0.589
5	4		1	0.582
2	1		1	0.573
6	5		1	0.558

```
(6 rows)
```